

## MapReduce & Apache Hadoop

The solutions to this exercise sheet should be demonstrated directly on the AWS cloud by using your own EC2 instance. You may either create a new EC2 instance for each member of your student group individually or share an instance among all members of your group. Please refer to our “GettingStarted-AWS” guide on Moodle for instructions on how to access your AWS account and create an EC2 instance. Note that you should always develop your code locally before uploading it to the AWS (refer also to our “GettingStarted” guide for instructions on setting up your local Hadoop environment).

The exercises once more require the `Wikipedia-EN-20120601_ARTICLES.tar.gz` and `Wikipedia-EN-20120601_KEYWORDS.TSV.gz` files which are provided for download on the course homepage on `moodle.uni.lu`. Upload these files from your local directory to your EC2 instance by issuing the following two commands:

- `scp -i <path-to-your/awskey.pem> </path-to-your/Wikipedia-EN-20120601_ARTICLES.tar.gz> <your-EC2-instance-DNS-name-or-IP-address>:~/`
- `scp -i <path-to-your/awskey.pem> </path-to-your/Wikipedia-EN-20120601_KEYWORDS.TSV.gz> <your-EC2-instance-DNS-name-or-IP-address>:~/`

### RATIO OF NOUNS

9 Points

#### Problem 1.

A sentence is composed of multiple words and each word has its specific grammatical category, also known as part-of-speech. Part-of-speech tagging (POS tagging) is the process of identifying the corresponding part-of-speech (verb, noun, adjective, etc.) of words in a sentence.

- From all words in the `Wikipedia-EN-20120601_ARTICLES.tar.gz` corpus, calculate the percentage of nouns in it by using two MapReduce jobs. **4 Points**
- Could you solve this exercise with a single job? Compare the running time of the solution using one iteration against the one using two. **5 Points**

*Hint: To identify the POS tags, you may use the OpenNLP library and a pre-trained POS tag model, which are available from `opennlp-tools-1.9.3.jar` and `opennlp-en-ud-ewt-pos-1.0-1.9.3.bin` on `moodle.uni.lu`. You may check the `PosTag` project on moodle to learn more on how to use the OpenNLP library.*

*Hint<sub>2</sub>: To chain multiple MapReduce jobs you may follow the following tutorial <https://towardsdatascience.com/chaining-multiple-mapreduce-jobs-with-hadoop-java-832a326cbfa7>.*

*Hint<sub>3</sub>: To execute external libraries on Hadoop - as the `opennlp-tools-1.9.3.jar` - you shall add the directory of those libraries to the `HADOOP_CLASSPATH` environment variable, as follows: `export HADOOP_CLASSPATH="PATH_TO_LIBRARIES/*"`*

**Exercise Sheet #2** DUE ON NOV. 18, 2021

---

INDEXING DOCUMENTS VIA HADOOP

12 Points

**Problem 2.**

1. Implement a pair of a **Map** and a **Reduce** function which, for each distinct term that occurs in any of the text documents in `Wikipedia-EN-20120601_ARTICLES.tar.gz`, counts the number of distinct documents in which the term appears. We will call this value the Document Frequency (DF) of that term in the entire set of Wikipedia articles. Store the resulting DF values of all terms in a single TSV file with the following schema:

`TERM<tab>DF`

*Hint: Also here consider the Porter stemmer that is available from `opennlp-tools-1.9.3.jar` on `moodle.uni.lu`. After importing the library with `import opennlp.tools.stemmer.PorterStemmer;` and creating an instance of `PorterStemmer` `stemmer = new PorterStemmer();` use `stemmer.stem(token);` for stemming an input token.*

6 Points

2. Implement another pair of a **Map** and a **Reduce** function which, for each document in `WikipediaEN-20120601_ARTICLES.tar.gz`, first counts the number of occurrences of each distinct term within the given document. We will call this value the Term Frequency (TF) of that term in the given document.

In a second step, your combined MapReduce function should multiply the TF value of each such term with the inverse of the logarithm of the normalized DF value calculated by the previous MapReduce function, i.e.,  $SCORE = TF \times \log(10000/DF + 1)$  for each combination of a term and a document. You may cache the former TSV file with the DF values by adding it via `Job.addCacheFile(<path-to-DF-file>)` in the driver function of your MapReduce program. The **Map** class should then load this file upon initialization into an appropriate main-memory data structure.

The result of this MapReduce function should be a single TSV file that has the same schema as the file we used in the previous exercise sheets:

6 Points

`ID<tab>TERM<tab>SCORE`

*Hint: You may use either the given document URL's or simple integer id's for the ID field of this TSV file (as long as they uniquely identify the original text article).*

**Exercise Sheet #2** DUE ON NOV. 18, 2021

---

PROCESSING JOINS VIA HADOOP

**9 Points**

**Problem 3.** For this exercise, consider either the file `Wikipedia-EN-20120601_KEYWORDS.TSV.gz` provided on `moodle.uni.lu` or your own TSV file you created for Problem 2 of this exercise sheet.

- Reconsider the first three queries of Problem 2 in Exercise Sheet #1 (called “Boolean Retrieval 1–3”).

This time, implement the three queries as pairs of **Map** and **Reduce** functions that perform *reduce-side joins* over the inverted lists for the four (stemmed) terms **rock**, **paper**, **scissors** and **game**.

Also compare your results (especially the runtime of the queries) with the ones you obtained earlier for PostgreSQL.

**3 Points for each query**